

Group Assignment 1: Breast Cancer Analysis using KNN

Aylin Kosar, Surya Aenuganti Ushasri, Salma Olmai, Nicolas Romero, Viraj Prasad Sapre

October 16, 2018

Upload Data

Surya:

```
#setwd("~/Fall 2018 R Data Files")
#Load the data from a csv file
cancer = read.csv("wisc_bc_data.csv", na.strings = NULL)
```

Check Data

Surya:

```
#Check the structure of the data
str(cancer)
```

```
## 'data.frame':    569 obs. of  32 variables:
## $ id              : int  87139402 8910251 905520 868871 9012568 906539 925291 87880 862989 89827 .
## $ diagnosis       : Factor w/ 2 levels "B","M": 1 1 1 1 1 1 1 2 1 1 ...
## $ radius_mean     : num  12.3 10.6 11 11.3 15.2 ...
## $ texture_mean    : num  12.4 18.9 16.8 13.4 13.2 ...
## $ perimeter_mean  : num  78.8 69.3 70.9 73 97.7 ...
## $ area_mean       : num  464 346 373 385 712 ...
## $ smoothness_mean : num  0.1028 0.0969 0.1077 0.1164 0.0796 ...
## $ compactness_mean : num  0.0698 0.1147 0.078 0.1136 0.0693 ...
## $ concavity_mean  : num  0.0399 0.0639 0.0305 0.0464 0.0339 ...
## $ points_mean     : num  0.037 0.0264 0.0248 0.048 0.0266 ...
## $ symmetry_mean   : num  0.196 0.192 0.171 0.177 0.172 ...
## $ dimension_mean  : num  0.0595 0.0649 0.0634 0.0607 0.0554 ...
## $ radius_se       : num  0.236 0.451 0.197 0.338 0.178 ...
## $ texture_se      : num  0.666 1.197 1.387 1.343 0.412 ...
## $ perimeter_se    : num  1.67 3.43 1.34 1.85 1.34 ...
## $ area_se         : num  17.4 27.1 13.5 26.3 17.7 ...
## $ smoothness_se   : num  0.00805 0.00747 0.00516 0.01127 0.00501 ...
## $ compactness_se  : num  0.0118 0.03581 0.00936 0.03498 0.01485 ...
## $ concavity_se    : num  0.0168 0.0335 0.0106 0.0219 0.0155 ...
## $ points_se       : num  0.01241 0.01365 0.00748 0.01965 0.00915 ...
## $ symmetry_se     : num  0.0192 0.035 0.0172 0.0158 0.0165 ...
## $ dimension_se    : num  0.00225 0.00332 0.0022 0.00344 0.00177 ...
## $ radius_worst    : num  13.5 11.9 12.4 11.9 16.2 ...
## $ texture_worst   : num  15.6 22.9 26.4 15.8 15.7 ...
## $ perimeter_worst : num  87 78.3 79.9 76.5 104.5 ...
## $ area_worst      : num  549 425 471 434 819 ...
## $ smoothness_worst : num  0.139 0.121 0.137 0.137 0.113 ...
## $ compactness_worst : num  0.127 0.252 0.148 0.182 0.174 ...
## $ concavity_worst : num  0.1242 0.1916 0.1067 0.0867 0.1362 ...
## $ points_worst    : num  0.0939 0.0793 0.0743 0.0861 0.0818 ...
## $ symmetry_worst  : num  0.283 0.294 0.3 0.21 0.249 ...
```

```
## $ dimension_worst : num 0.0677 0.0759 0.0788 0.0678 0.0677 ...
```

Organize Data

Surya:

```
table(cancer$diagnosis)
```

```
##  
## B M  
## 357 212
```

```
#Check whether there is any missing data  
sum(is.na(cancer$diagnosis))
```

```
## [1] 0
```

```
#Clearly this should be a factor hence converting it to a factor and labeling the levels to benign or malignant  
cancer$diagnosis <- factor(cancer$diagnosis, levels = c("B", "M"), labels = c("benign", "malignant"))  
levels(cancer$diagnosis)
```

```
## [1] "benign" "malignant"
```

Surya: Since we need to measure distances to classify them according to knn, we need the variables to have numerical values on a same scale. So we normalize the variables. Here we are creating a function so that we can apply the normalization to all columns using this single function.

```
normalize = function(x) {  
  y = (x - min(x))/(max(x) - min(x))  
  y  
}
```

```
#Applying the above normalization function to all columns except the first 2  
#so lapply is a function in r where can specify the function to apply and the columns on which we have  
can_n_L <- lapply(cancer[, 3:32], normalize)  
#converting the data to a data frame[Since wdbc_n_L consist of only data from #3- 32 columns]  
can_n <- data.frame(can_n_L)
```

```
can_n[1:3, 1:4]
```

```
## radius_mean texture_mean perimeter_mean area_mean  
## 1 0.2526859 0.0906324 0.2422777 0.13599152  
## 2 0.1712812 0.3124789 0.1761454 0.08606575  
## 3 0.1921056 0.2407846 0.1874784 0.09743372
```

```
rownames(can_n) <- cancer$id
```

```
#Isolate the class labels and name them accordingly
```

```
BM_class <- cancer[, 2]
```

```
#names(BM_class)-> this would give null because there are no labels yet because #there #are no attributes
```

```
names(BM_class) <- cancer$id
```

```
BM_class[1:3]
```

```
## 87139402 8910251 905520  
## benign benign benign  
## Levels: benign malignant
```

```
#so now each label comes under an attribute which is actually the row name/id  
#imagine a single row but 569 of attributes
```

Creating training and test (validation) datasets

Surya:

```
nrow(cancer)

## [1] 569

rand_permute <- sample(x = 1:569, size = 569)

rand_permute[1:5]

## [1] 500 164 100 168 414

# save(rand_permute, file='rand_permute.RData')

#load("rand_permute.RData")

#randomly permute the rows of data
all_id_random = cancer[rand_permute, "id"]

# Select the first third of these for validation

569/3

## [1] 189.6667

#Get the first 1/3 ids of the data and keep it for validation
validate_id <- as.character(all_id_random[1:189])
#Get the next 2/3 ids of the data and keep it for training
training_id <- as.character(all_id_random[190:569])
```

Surya: Subset the data by taking the data of the respective ids

```
can_train <- can_n[training_id, ]

can_val <- can_n[validate_id, ]

BM_class_train <- BM_class[training_id]

BM_class_val <- BM_class[validate_id]

table(BM_class_train)

## BM_class_train
##    benign malignant
##      240      140

table(BM_class_val)

## BM_class_val
##    benign malignant
##      117      72
```

Executing knn

Surya:

```
library(class)
```

```

`?`(knn)

## starting httpd help server ... done
sqrt(nrow(can_train))

## [1] 19.49359
k = 19

#Fitting the model and validating against test set
knn_predict = knn(can_train, can_val, BM_class_train, k = 19)

knn_predict[1:3]

## [1] malignant benign    malignant
## Levels: benign malignant
#Check the confusion matrix for true positives and true negatives
table(knn_predict, BM_class_val)

##           BM_class_val
## knn_predict benign malignant
##   benign      116         7
##   malignant     1        65
prop.table(table(knn_predict, BM_class_val))

##           BM_class_val
## knn_predict      benign    malignant
##   benign  0.613756614 0.037037037
##   malignant 0.005291005 0.343915344

```

It really depends on the randomly selected data for testing and validating

Testing different values of k

Aylin: The knn numerical value are given random variables to predict the outcome of the train set. The first is considered the best.

```

knn_predict_3 = knn(can_train, can_val, BM_class_train, k = 3)

knn_predict_7 = knn(can_train, can_val, BM_class_train, k = 7)

knn_predict_11 = knn(can_train, can_val, BM_class_train, k = 11)

knn_predict_31 = knn(can_train, can_val, BM_class_train, k = 31)

table(knn_predict_3, BM_class_val)

##           BM_class_val
## knn_predict_3 benign malignant
##   benign      114         3
##   malignant     3        69
table(knn_predict_7, BM_class_val)

##           BM_class_val
## knn_predict_7 benign malignant
##   benign      115         2

```

```
##      malignant      2      70
table(knn_predict_11, BM_class_val)
```

```
##              BM_class_val
## knn_predict_11 benign malignant
##      benign      116      4
##      malignant      1      68
table(knn_predict_31, BM_class_val)
```

```
##              BM_class_val
## knn_predict_31 benign malignant
##      benign      116      8
##      malignant      1      64
```

Study significance of the variables

Aylin: Below the names of the data set are listed. A model is binomial regression model is created. The original name of the model was `lm_1` but I renamed it `g1` to make it easier. “can” is also substituted for “wbcd”. The names of the model is created in the second code after the `g1` model is created. Then the F statistic is created.

```
names(can_train)
```

```
## [1] "radius_mean"      "texture_mean"      "perimeter_mean"
## [4] "area_mean"        "smoothness_mean"   "compactness_mean"
## [7] "concavity_mean"    "points_mean"       "symmetry_mean"
## [10] "dimension_mean"    "radius_se"         "texture_se"
## [13] "perimeter_se"      "area_se"           "smoothness_se"
## [16] "compactness_se"    "concavity_se"      "points_se"
## [19] "symmetry_se"       "dimension_se"       "radius_worst"
## [22] "texture_worst"     "perimeter_worst"   "area_worst"
## [25] "smoothness_worst"  "compactness_worst" "concavity_worst"
## [28] "points_worst"      "symmetry_worst"    "dimension_worst"
```

```
g1 = lm(radius_mean ~ BM_class_train, data = can_train)
summary(g1)
```

```
##
## Call:
## lm(formula = radius_mean ~ BM_class_train, data = can_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.30429 -0.07569  0.00025  0.07136  0.50786
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.245812   0.007394   33.24  <2e-16 ***
## BM_class_trainmalignant 0.246324   0.012182   20.22  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1146 on 378 degrees of freedom
## Multiple R-squared:  0.5196, Adjusted R-squared:  0.5183
## F-statistic: 408.9 on 1 and 378 DF, p-value: < 2.2e-16
```

```
names(summary(g1))
```

```
## [1] "call"          "terms"          "residuals"      "coefficients"
## [5] "aliased"        "sigma"          "df"             "r.squared"
## [9] "adj.r.squared" "fstatistic"     "cov.unscaled"
```

```
summary(g1)$fstatistic
```

```
## value numdf dendf
## 408.8577 1.0000 378.0000
```

```
# The significance measure we want:
```

```
summary(g1)$fstatistic[1]
```

```
## value
## 408.8577
```

Aylin: The first chunk of code, a vector is created in order to keep all the outputs together. The next code the variables are run through to try to get a linear fit and have the F- statistic stored. The first code also asks NA to be repeated 30 times. The first three variables in the first row has the f statistic value.

```
exp_var_fstat <- as.numeric(rep(NA, times = 30))
```

```
names(exp_var_fstat) <- names(can_train)
```

```
exp_var_fstat["radius_mean"] <- summary(lm(radius_mean ~ BM_class_train, data = can_train))$fstatistic[1]
```

```
exp_var_fstat["texture_mean"] <- summary(lm(texture_mean ~ BM_class_train, data = can_train))$fstatistic[1]
```

```
exp_var_fstat["perimeter_mean"] <- summary(lm(perimeter_mean ~ BM_class_train, data = can_train))$fstatistic[1]
```

```
exp_var_fstat
```

```
## radius_mean texture_mean perimeter_mean area_mean
## 408.85767 78.54968 444.85665 NA
## smoothness_mean compactness_mean concavity_mean points_mean
## NA NA NA NA
## symmetry_mean dimension_mean radius_se texture_se
## NA NA NA NA
## perimeter_se area_se smoothness_se compactness_se
## NA NA NA NA
## concavity_se points_se symmetry_se dimension_se
## NA NA NA NA
## radius_worst texture_worst perimeter_worst area_worst
## NA NA NA NA
## smoothness_worst compactness_worst concavity_worst points_worst
## NA NA NA NA
## symmetry_worst dimension_worst
## NA NA
```

Looping over variable names

Aylin: The last step is repeated again to create a vector in order to hold the significance measures. The code commented out produces an error since there is no variable with the name form exp_vars[j]. The named variable needs to be stored in the variable named “slot”, so a variable is created and a formula is created in order for this to happen in the next code snippet below the code with the error. Again the variable exp_var_fstat is called and a table is outputted with the variables in the data set and the f statistic.

```

exp_vars = names(can_train)

exp_var_fstat = as.numeric(rep(NA, times = 30))

names(exp_var_fstat) = exp_vars

# Code snippet commented out creates an error.

#for (j in 1:length(exp_vars)) {
#  exp_var_fstat[exp_vars[j]] = summary(lm(exp_vars[j] ~ BM_class_train, data = can_train))$fstatist
# }

for (j in 1:length(exp_vars)) {

  exp_var_fstat[exp_vars[j]] =

    summary(lm(as.formula(paste(exp_vars[j], " ~ BM_class_train")), data = can_train))$fstatistic[1]

}

exp_var_fstat

```

##	radius_mean	texture_mean	perimeter_mean	area_mean
##	4.088577e+02	7.854968e+01	4.448566e+02	3.523234e+02
##	smoothness_mean	compactness_mean	concavity_mean	points_mean
##	5.206519e+01	2.024908e+02	4.308096e+02	6.096495e+02
##	symmetry_mean	dimension_mean	radius_se	texture_se
##	5.392311e+01	2.317485e-02	1.704814e+02	4.948036e-03
##	perimeter_se	area_se	smoothness_se	compactness_se
##	1.575166e+02	1.436926e+02	1.198831e+00	3.769899e+01
##	concavity_se	points_se	symmetry_se	dimension_se
##	5.853678e+01	9.741701e+01	5.988748e-02	2.060474e+00
##	radius_worst	texture_worst	perimeter_worst	area_worst
##	5.567013e+02	1.037929e+02	5.907298e+02	4.119974e+02
##	smoothness_worst	compactness_worst	concavity_worst	points_worst
##	9.010266e+01	1.993949e+02	3.540642e+02	6.870139e+02
##	symmetry_worst	dimension_worst		
##	9.029549e+01	4.150027e+01		

Aylin: The function `lapply` or `sapply` is used to avoid initializing the variables. This gets all stored in a second variable “exp_var_fstat2”.

```

exp_var_fstat2 = sapply(exp_vars, function(x) {

  summary(lm(as.formula(paste(x, " ~ BM_class_train")), data = can_train))$fstatistic[1]

})

exp_var_fstat2

```

##	radius_mean.value	texture_mean.value	perimeter_mean.value
##	4.088577e+02	7.854968e+01	4.448566e+02
##	area_mean.value	smoothness_mean.value	compactness_mean.value
##	3.523234e+02	5.206519e+01	2.024908e+02
##	concavity_mean.value	points_mean.value	symmetry_mean.value

```
##          4.308096e+02          6.096495e+02          5.392311e+01
## dimension_mean.value      radius_se.value      texture_se.value
##          2.317485e-02          1.704814e+02          4.948036e-03
##      perimeter_se.value      area_se.value      smoothness_se.value
##          1.575166e+02          1.436926e+02          1.198831e+00
## compactness_se.value      concavity_se.value      points_se.value
##          3.769899e+01          5.853678e+01          9.741701e+01
##      symmetry_se.value      dimension_se.value      radius_worst.value
##          5.988748e-02          2.060474e+00          5.567013e+02
## texture_worst.value      perimeter_worst.value      area_worst.value
##          1.037929e+02          5.907298e+02          4.119974e+02
## smoothness_worst.value      compactness_worst.value      concavity_worst.value
##          9.010266e+01          1.993949e+02          3.540642e+02
##      points_worst.value      symmetry_worst.value      dimension_worst.value
##          6.870139e+02          9.029549e+01          4.150027e+01
```

```
names(exp_var_fstat2) = exp_vars
```

plyr version of the fit

Aylin: The data is now combined together by creating a list of data.frames with one category for each variable. The BM class variable is packaged into the data.frames so all the variables are all in one location. When you get the output you will see numerical values with four categories:sample, variable, value, and the variable's class.

```
can_L = lapply(exp_vars, function(x) {
  df = data.frame(sample = rownames(can_train), variable = x, value = can_train[,
    x], class = BM_class_train)
  df
})

head(can_L[[1]])
```

```
##      sample      variable      value      class
## 874373      874373 radius_mean 0.2238156      benign
## 8510653     8510653 radius_mean 0.2886554      benign
## 864033      864033 radius_mean 0.1323300      benign
## 924632      924632 radius_mean 0.2791897      benign
## 874217      874217 radius_mean 0.5361825      malignant
## 901034302   901034302 radius_mean 0.2630981      benign
```

```
head(can_L[[5]])
```

```
##      sample      variable      value      class
## 874373      874373 smoothness_mean 0.4072402      benign
## 8510653     8510653 smoothness_mean 0.4953507      benign
## 864033      864033 smoothness_mean 0.4610454      benign
## 924632      924632 smoothness_mean 0.2581927      benign
## 874217      874217 smoothness_mean 0.3001715      malignant
## 901034302   901034302 smoothness_mean 0.1961722      benign
```

```
names(can_L) = exp_vars
```

Aylin: The function lapply in the plyr library. The function sapply can also be the same since they are basically the same function. There are three different types of mean(radius_mean, texture_mean, perimeter_mean) along with the f statistic values created.


```
library(plyr)

var_sig_fstats = laply(can_L, function(df) {
  fit = lm(value ~ class, data = df)
  f = summary(fit)$fstatistic[1]
  f
})

names(var_sig_fstats) = names(can_L)

var_sig_fstats[1:3]

##      radius_mean      texture_mean perimeter_mean
##      408.85767       78.54968      444.85665
```

Conclusions about significance of the variables

Aylin: The first code snippet is asking for the data for variables ordered 1 to 5 which is points_worst, perimeter_worst, points_mean, radius_worst, and area_worst. It then prints out for each of these variables the significant f stats for each. The same goes for the second code snippet except for data variables ordered 25 to 30. Below, the last code snippet, the variables in the training set named data.frame are reordered by significance in order to prepare to do the kNN.

```
most_sig_stats = sort(var_sig_fstats, decreasing = T)

most_sig_stats[1:5]

##      points_worst      points_mean perimeter_worst      radius_worst
##      687.0139       609.6495      590.7298      556.7013
##      perimeter_mean
##      444.8566

most_sig_stats[25:30]

## compactness_se      dimension_se      smoothness_se      symmetry_se      dimension_mean
##      37.698991861      2.060474129      1.198830786      0.059887480      0.023174852
##      texture_se
##      0.004948036

can_train_ord = can_train[, names(most_sig_stats)]
```

Monte Carlo Cross-Validation

Selection of the family of training sets

Aylin: The data below is subsetting. The first code takes the length of the training set, the second takes the length of the training set and multiplies it by 2/3, the third takes the length of the training set and subtracts it from 253. The value 253 is the size of the new training set. The training data gets loaded and is named training_family_L.Data.

```
length(training_id)

## [1] 380

(2/3) * length(training_id)

## [1] 253.3333
```

```
length(training_id) - 253

## [1] 127
# Use 253 as the training set size.

training_family_L = lapply(1:1000, function(j) {
  perm = sample(1:380, size = 380, replace = F)
  shuffle = training_id[perm]
  trn = shuffle[1:253]
  trn
})

# save(training_family_L, file='training_family_L.RData')

#load("training_family_L.RData")

validation_family_L = lapply(training_family_L, function(x) setdiff(training_id, x))
```

Finding an optimal set of variables and optimal k

Aylin: The code below calculates the distributions of errors over the 1000 training-validation pairs for different subsets of the variables. The square root of the reference set size is taken in order to test options for k. The value will vary from 3 to 19 and the last code, for each training - validation subset, number of variables, and k, the error of the kNN prediction in the validation set is calculated.

```
N = seq(from = 3, to = 29, by = 2)

sqrt(length(training_family_L[[1]]))

## [1] 15.90597

K = seq(from = 3, to = 19, by = 2)

1000 * length(N) * length(K)
```

```
## [1] 126000
```

Execution of the test with loops

Aylin: The data frame will be initialized to store 126,000 entries. A new function for the core kNN error is created and stored in the knn_test, n = 5 and k = 7.

```
paramter_errors_df = data.frame(mc_index = as.integer(rep(NA, times = 126000)),
  var_num = as.integer(rep(NA, times = 126000)), k = as.integer(rep(NA, times = 126000)),
  error = as.numeric(rep(NA, times = 126000)))

knn_test = knn(train = can_train_ord[training_family_L[[1]], 1:5], test = can_train_ord[validation_fami
  1:5], cl = BM_class_train[training_family_L[[1]]], k = 7)

knn_test[1:3]

## [1] benign benign benign
## Levels: benign malignant

tbl_test = table(knn_test, BM_class_train[validation_family_L[[1]])
```

```
tbl_test
```

```
##  
## knn_test      benign malignant  
##    benign      80          3  
##    malignant   1          43
```

```
err_rate = (tbl_test[1, 2] + tbl_test[2,1])/length(validation_family_L[[1]])
```

```
err_rate
```

```
## [1] 0.03149606
```

Aylin: Another function is created to run the code snippet and return back the error rate along with other parameters. Then below after “sample” commented out, a nested for loop is created.

```
# j = index, n = length of range of variables, k=k
```

```
core_knn = function(j, n, k) {  
  knn_predict = knn(train = can_train_ord[training_family_L[[j]], 1:n],  
                    test = can_train_ord[validation_family_L[[j]], 1:n], cl = BM_class_train[training_family_L[[j]])  
  
  tbl = table(knn_predict, BM_class_train[validation_family_L[[j]]])  
  err = (tbl[1, 2] + tbl[2, 1])/length(validation_family_L[[j]])  
  err  
}
```

```
# sample
```

```
core_knn(1, 5, 7)
```

```
## [1] 0.03149606
```

```
iter = 1
```

```
str_time = Sys.time()
```

```
for (j in 1:1000) {  
  for (n in 1:length(N)) {  
    for (m in 1:length(K)) {  
      err = core_knn(j, N[n], K[m])  
      paramter_errors_df[iter, ] <- c(j, N[n], K[m], err)  
      iter = iter + 1  
    }  
  }  
}
```

```
time_lapsed_for = Sys.time() - str_time
```

```
save(paramter_errors_df, time_lapsed_for, file = "for_loop_paramter_errors.RData")
```

```
load("for_loop_paramter_errors.RData")
```

```
time_lapsed_for
```

```
## Time difference of 8.082501 mins
```

Execution with plyr

Aylin: A data frame of all possible parameter combinations is created. Then they will be nested inside the several `**ply` functions. Below is just a trst using 20 choices of parameters and then you do a full run.

```
param_df1 = merge(data.frame(mc_index = 1:1000), data.frame(var_num = N))

param_df = merge(param_df1, data.frame(k = K))

str(param_df)

## 'data.frame':    126000 obs. of  3 variables:
## $ mc_index: int  1 2 3 4 5 6 7 8 9 10 ...
## $ var_num : num  3 3 3 3 3 3 3 3 3 3 ...
## $ k       : num  3 3 3 3 3 3 3 3 3 3 ...

knn_err_est_df_test_test = ddply(param_df[1:20, ], .(mc_index, var_num, k), function(df) {

  err = core_knn(df$mc_index[1], df$var_num[1], df$k[1])

  err

})

head(knn_err_est_df_test_test)

##   mc_index var_num k          V1
## 1         1      3 3 0.05511811
## 2         2      3 3 0.07874016
## 3         3      3 3 0.04724409
## 4         4      3 3 0.09448819
## 5         5      3 3 0.07086614
## 6         6      3 3 0.07874016

str_time = Sys.time()

knn_err_est_df_test_test = ddply(param_df, .(mc_index, var_num, k), function(df) {
  err = core_knn(df$mc_index[1], df$var_num[1], df$k[1])

  err

})

time_lapsed = Sys.time() - str_time

save(knn_err_est_df_test_test, time_lapsed, file = "knn_err_est_df_test_test")

load("knn_err_est_df_test_test")

time_lapsed

## Time difference of 4.032573 mins

head(knn_err_est_df_test_test)

##   mc_index var_num k          V1
## 1         1      3 3 0.05511811
## 2         1      3 5 0.05511811
## 3         1      3 7 0.05511811
## 4         1      3 9 0.03937008
## 5         1      3 11 0.03149606
```

```
## 6      1      3 13 0.03149606
names(knn_err_est_df_test_test)[4] = "error"
```

Getting summary performance statistics

Aylin: The mean is taken for each of the parameters. Then the mean error is taken for the parameters.

```
mean_ex_df = subset(knn_err_est_df_test_test, var_num == 5 & k == 7)

head(mean_ex_df)
```

```
##      mc_index var_num k      error
## 12          1      5 7 0.03149606
## 138         2      5 7 0.06299213
## 264         3      5 7 0.03937008
## 390         4      5 7 0.06299213
## 516         5      5 7 0.04724409
## 642         6      5 7 0.07874016
```

```
mean(mean_ex_df$error)
```

```
## [1] 0.05548819
```

```
mean_errs_df = dplyr::group_by(knn_err_est_df_test_test, var_num, k) %>%
  summarise(mean_error = mean(error))

head(mean_errs_df)
```

```
##   var_num k      V1
## 1      3 3 0.06101575
## 2      3 5 0.06074803
## 3      3 7 0.06195276
## 4      3 9 0.05985827
## 5      3 11 0.05826772
## 6      3 13 0.05629921
```


```
names(mean_errs_df)[3] = "mean_error"
```

Survey of parameter performance


Aylin: The performance gets visualized using the library ggplot.

```
library(ggplot2)
```

```
ggplot(data = mean_errs_df, aes(x = var_num, y = k, color = mean_error)) + geom_point(size = 10) +
```



```
ggplot(data = subset(mean_errs_df, var_num >= 15), aes(x = var_num, y = k, color = mean_error)) +  
  geom_point(size = 10) + theme_bw()
```



Aylin: Variables with a low k works the best.

```
subset(mean_errs_df, var_num == 17)
```

```
##   var_num  k mean_error
## 64     17  3 0.04700787
## 65     17  5 0.04331496
## 66     17  7 0.04236220
## 67     17  9 0.04225984
## 68     17 11 0.04317323
## 69     17 13 0.04371654
## 70     17 15 0.04346457
## 71     17 17 0.04278740
## 72     17 19 0.04233858
```

```
subset(mean_errs_df, var_num == 19)
```

```
##   var_num  k mean_error
## 73     19  3 0.03497638
## 74     19  5 0.03340157
## 75     19  7 0.03388189
## 76     19  9 0.03425197
## 77     19 11 0.03548819
## 78     19 13 0.03650394
## 79     19 15 0.03722835
## 80     19 17 0.03821260
## 81     19 19 0.03914173
```

```
subset(mean_errs_df, var_num == 21)
```

```
##      var_num k mean_error
## 82        21 3 0.03916535
## 83        21 5 0.03667717
## 84        21 7 0.03436220
## 85        21 9 0.03422835
## 86        21 11 0.03485827
## 87        21 13 0.03515748
## 88        21 15 0.03604724
## 89        21 17 0.03710236
## 90        21 19 0.03833071
```

```
subset(mean_errs_df, var_num == 25)
```

```
##      var_num k mean_error
## 100        25 3 0.03818898
## 101        25 5 0.03507874
## 102        25 7 0.03570866
## 103        25 9 0.03469291
## 104        25 11 0.03465354
## 105        25 13 0.03588189
## 106        25 15 0.03707874
## 107        25 17 0.03825197
## 108        25 19 0.03969291
```

```
mean_errs_df[which.min(mean_errs_df$mean_error), ]
```

```
##      var_num k mean_error
## 74         19 5 0.03340157
```

```
names(can_train_ord)
```

```
## [1] "points_worst"      "points_mean"       "perimeter_worst"
## [4] "radius_worst"      "perimeter_mean"    "concavity_mean"
## [7] "area_worst"        "radius_mean"       "concavity_worst"
## [10] "area_mean"         "compactness_mean"  "compactness_worst"
## [13] "radius_se"         "perimeter_se"      "area_se"
## [16] "texture_worst"     "points_se"         "symmetry_worst"
## [19] "smoothness_worst"  "texture_mean"      "concavity_se"
## [22] "symmetry_mean"     "smoothness_mean"   "dimension_worst"
## [25] "compactness_se"    "dimension_se"      "smoothness_se"
## [28] "symmetry_se"       "dimension_mean"    "texture_se"
```

Validation of the final test

VIRAJ

```
bcd_val_ord = can_val[, names(can_train_ord)] #Here the 189 observations are taken from both the tables
```

```
bm_val_pred <- knn(train = can_train_ord[, 1:27], can_val[, 1:27], BM_class_train,
  k = 3) #training both these variables with k=3 for the class train table
```

```
tbl_bm_val <- table(bm_val_pred, BM_class_val) # pred table contains the b or m predicated values, clas
tbl_bm_val
```

```
##      BM_class_val
```



```
## bm_val_pred benign malignant
##   benign      113      37
##   malignant    4      35
(val_error <- tbl_bm_val[1, 2] + tbl_bm_val[2, 1])/length(BM_class_val) #putting these values from tbl_
## [1] 0.2169312
```